

PROFESSIONAL SUMMARY

Data Engineer with 5+ years of experience designing and optimizing scalable data pipelines using PySpark, SQL, and Python across AWS and GCP environments. Proven expertise in cloud data lakes, lakehouse architectures, and analytics-ready data models supporting large-scale business reporting.

CORE SKILLS

- **Programming & Scripting:** Python, PySpark, SQL, Shell Scripting, Java
- **Big Data & Precessing:** Apache Spark, Hive, Apache Iceberg
- **Cloud Platforms & Data Platforms:** AWS (Glue, S3, Athena, Lambda, Redshift, SQS, SNS), GCP (BigQuery, Cloud Storage, GCP Cloud Composer)
- **Data Formats:** Parquet, Avro, ORC
- **Orchestration & Automation:** Apache Airflow, AWS Batch, lambda, TWS Scheduler
- **Version Control:** Git (GitHub, GitLab)

PROFESSIONAL EXPERIENCE

Deloitte Touche Tohmatsu Ltd

India

Consultant – Data Engineering and AI

Sep 2025 – Present

- Data Engineer for Eli Lilly's IMPI platform, designing and maintaining near-real-time simulation pipelines supporting clinical inventory optimization and patient allocation decisions, resulting in savings of over 30 man-hours per week and reducing manual effort for a team of 23 personnel.
- Built hourly simulation pipelines delivering near real-time availability insights across global clinical sites, improving planning accuracy and operational responsiveness.
- Engineered scalable real-time and batch data processing architectures using AWS Glue, Apache Flink, PySpark, Python, Lambda, and Apache Airflow for high-volume pharmaceutical datasets.
- Developed resilient ELT workflows using AWS Glue, Airflow, and PySpark to support distributed data processing and downstream analytics.
- Optimized data storage and retrieval using Amazon S3 for staging and Aurora PostgreSQL for structured reference data, analytical workloads, and inventory reporting.
- Implemented automated and scalable processing pipelines using AWS Batch, Lambda, and Fargate-based containerized services to improve reliability and operational efficiency.
- Designed complex SQL transformations and analytical queries within Aurora PostgreSQL to ensure data consistency across reporting platforms and simulation systems.
- Built and maintained cloud-native orchestration and monitoring solutions to improve observability, fault tolerance, and pipeline resiliency using Grafana
- Configured DMS jobs to move data from on prem Oracle database to RDS(Aurora Postgres)
- Collaborated with cross-functional global teams in US , Egypt and Ireland in agile environments to deliver scalable and high-performance data engineering solutions.

Capgemini SE

India

Consultant – Data Engineering & Analytics

Aug 2024 – Aug 2025

- Worked Lakehouse migration using Apache Iceberg for Global UK Bank, reducing data reprocessing effort by 30 percent and improved consistency
- Developed Python-driven transformation logic supporting SCD Type 1 & Type 2, enabling accurate historical tracking for downstream analytics
- Migrated legacy data lake architecture to a Lakehouse model using Apache Iceberg, improving data consistency and reducing reprocessing.
- Orchestrated workflows using Apache Airflow (Python DAGs), achieving 90% pipeline reliability.
- Wrote and optimized SQL queries in Athena and Redshift, reducing reporting latency for analytics teams
- Implemented logging and monitoring via CloudWatch.

- Contributed to a Data Hydration initiative by creating data mesh instances for datasets stored in Amazon S3, enhancing data discoverability and reusability across the organization with enhanced data governance

Cognizant

India

Associate – Data Engineer

Sep 2021 – Jul 2024

- Developed end-to-end ETL pipelines using PySpark, SQL, and Shell scripting for a US banking client, supporting enterprise MDM systems.
- Built a custom ingestion framework with integrated data quality checks, improving data accuracy .
- Standardized and transformed large datasets into analytics-ready data marts, enabling faster consumption by business teams.
- Automated legacy ETL workflows, reducing manual intervention and improving pipeline execution%.
- Migrated Hadoop/Hive datasets to cloud storage (GCS) and converted data formats to Avro, improving scalability and reducing infrastructure costs.
- Designed analytical datasets and supported dashboards in Looker Studio, improving business reporting turnaround time.
- Worked closely with stakeholders for pre-acceptance testing and validation, ensuring zero production rework.

EDUCATION

University of Visvesvaraya College of Engineering (UVCE)

Bachelor of Technology in Engineering

Jul 2017 – Sep 2021

CERTIFICATIONS

- AWS Certified Data Engineer Associate: DEA-C01
- AWS Certified AI Practitioner: AIF-C01
- Associate Data Practitioner - Google Cloud
- IBM Data Engineering: IBM, Coursera
- Machine Learning: Stanford University, Coursera
- Microsoft Certified: Azure Fundamentals